


AG DSN NAT Gateway Concept

Felix Kluge

Simon Hanisch

February 28, 2017



AG DSN
AG Dresdner Studentennetz

Contents

1	Introduction	3
1.1	Motivation	3
1.2	Nomenclature	3
1.3	Features	4
1.3.1	Private Networks	4
1.3.2	Roaming Networks	4
1.3.3	Rate Limiting	4
1.3.4	Failover	4
1.3.5	Port Forwarding	5
1.3.6	Internal API	5
2	Basics	6
2.1	IP Mobility	6
2.2	Netfilter Hooks	6
3	Management Component	8
3.1	Structure	8
3.2	Data Model	8
3.2.1	Translations	8
3.2.2	Forwardings	8
3.2.3	Throttles	8
3.3	Program State	8
3.4	State Updates	9
4	Implementation	10
4.1	packet filter	10
4.2	pfSense	10
4.3	Iptables	10
4.4	Nftables	10
4.5	Conntrack / Conntrackd	11
4.6	Keepalived	12
5	Setup	13
5.1	Live-Setup	13
5.2	Gateways	13
6	Evaluation	16
6.1	Introduction	16
6.2	Setup	16
6.3	Generating Traffic	16
6.4	Metrics	17

6.5	Problems	17
6.6	Results	18
7	Appendix	19
7.1	keepalived.conf	19
7.2	nping	20
7.2.1	Example usage	20
7.2.2	Python script	20
7.3	nftables	21
7.3.1	map	21
7.3.2	tree	21
7.3.3	HowTo	22
7.4	iptables	23
7.4.1	tree	23
7.4.2	HowTo	23
7.4.3	conntrackd.conf	24
7.5	Application config	25

Chapter 1

Introduction

1.1 Motivation

The network connections in the local residential homes currently suffer from a simple but substantial problem: only IPv4 addresses are provided by the university and there are only enough addresses to assign one to each student. The fact that everybody owns more than one IP address capable device today leads to the need of network address translation (NAT). The situation right now is that every student needs to buy a small router that performs the NAT for him. This solution works in the current setup but interferes with our plans to provide infrastructure WLAN for our users. With infrastructure WLAN the user no longer connects to his own private network, meaning that his router can no longer perform the NAT for him. The NAT has to be provided by the AG DSN. Another side effect of this is that the user no longer needs to buy a router, which leads to less wireless networks interfering with our infrastructure WLAN. The goal of this project is to build a central NAT solution for all our users.

1.2 Nomenclature

Network Address Translation (NAT) Network Address Translation is the process of exchanging the address information of a network packet. It is often required, because there are not enough public IPv4 addresses for all hosts of a network. A private network is used for those hosts and the internal IP addresses get translated by a NAT gateway before the packet leaves the network. The gateway tracks the connection of each host in the network to the public network and thereby the internal host that needs to receive the related incoming traffic from the public network.

Source NAT (SNAT) The source IP address gets exchanged by the gateway, this is the normal NAT mode in use by consumer routers.

Destination NAT (DNAT) Destination NAT is required if a service located on an internal host should be available on the public network. When the gateway receives traffic from the public network on a specific port the traffic is forwarded to the host (usually known as port forwarding in customer routers). So the destination address of a packet gets changed without the internal host creating a connection first.

Classless Inter-Domain Routing (CIDR) The term Classless Inter-Domain Routing describes an approach to optimize the usage of the 32 bit IPv4 address space. It is used in this document to describe subnets and their size.

1.3 Features

1.3.1 Private Networks

Every user of the network gets a private network (also referred to as “home network”). The term “network” refers to a VLAN according to IEEE 802.1Q[?]. This network will be located at the router next to the users flat. If the user tries to connect to the network via LAN or WLAN inside his building, he will be placed inside this network. This implies that all entirety networks are only distinct inside the building, not in a global (whole of residential homes) context. The IPs assigned inside a private network are going to be RFC6598 addresses[?], in practice an /24 network out of the Carrier-Grade NAT (CGN) range (100.64.0.0/10). Every private network will be mapped via network address translation (NAT[?]) to an external IP address on the NAT gateway. Every external IP address is only in use by a single user.

1.3.2 Roaming Networks

All buildings get a unique network (VLAN) for roaming purposes. When a user connects to the network at another building, i.e. not where his flat is located, he is placed inside this local roaming network. These networks will contain an subnet from the CGN allocation. Per default, connections from roaming networks are not translated to external IP addresses. As a result, no connectivity to addresses outside of the internal network is possible. During the authorization of a client in an roaming network, a rule is created on the NAT gateway. This rule translates connections to the assigned private IP address to the public IP address of the user. Another approach to solve these roaming issues is IP mobility as described in RFC3344[?]. It is described in the chapter IP Mobility 2.1.

1.3.3 Rate Limiting

Due to an agreement with the university, the network connection for individual users must be limited if a certain amount of traffic is exceeded. Instead of blocking the access to the network for those users, rate limiting will be performed by the NAT gateway. It is possible to configure exceptions, for certain destinations, that are excluded from the rate limiting.

1.3.4 Failover

The setup allows multiple NAT gateways to provide high availability. If one gateway fails, the network stays operational. To achieve failover, the connection state needs to be synchronized. The tool to synchronize the connection tracking state between multiple hosts running linux (see Setup 5.1 chapter for details) is conntrackd[?]. Besides the connection tracking state, the availability of the gateways is ensured by using keepalived[?], a VRRP daemon for Linux. For failover, multiple setups are possible.

Active - Passive Only one gateway is active, the other one is passive and becomes active if the other gateway fails. This setup is easy to set up but wastes resources, because one gateway is idling.

Active - Active (Synchronous Routing) The internal hosts are split up into two sets and each gateway gets to perform NAT for one set. So both gateways are active and can be used as backup for the other gateway. The load balancing in this case is static and can be performed with policy based routing. Both directions of a connection get translated by the same gateway (Synchronous Routing).

Active - Active (Asynchronous Routing) This is the most complex setup, because it is not clear which gateway performs the translation. This is why the connection tracking state has to be synchronized continuously and with a very low latency. Example: A request by an internal host is translated by gateway X and the response is translated by gateway Y. During the time the response needs to arrive, the connection tracking state has to be synchronized – otherwise, the packets are rejected. This scenario allows dynamic load balancing between multiple NAT gateways without knowledge which gateway a packet is processed on.

1.3.5 Port Forwarding

NAT prevents users to use remotely initiated connections. To address this issue, port forwardings are configurable. In detail, the NAT gateway allows incoming connections to the external (“public”) IP addresses and forward them to an configured internal host inside a private network. Related to port forwarding, a feature set comparable to publicly sold routers is available. Popular examples are TP-Link routers (e.g. TL-WR841ND[?]) and AVM routers (e.g. Fritzbox 4020[?]), which allow the user to configure forwarding rules. Criteria for these rules are the inbound (external) port and the protocol. The redirection is possible to an configurable internal IP address and port.

1.3.6 Internal API

To set up the NAT gateway, an interface to the configuration tools already in place is required. This targets the user management tool and the user self-service tool. In addition, the DHCP server for roaming networks needs to be able to perform updates on the NAT gateway. These other parts of the infrastructure are not part of this project, but the interface of the gateway should be designed for their needs.

Chapter 2

Basics

2.1 IP Mobility

IP mobility[?][?] uses a complex setup to allow roaming users to retain an IP address from their home network. An endpoint in the home network of the user - called the home agent - forwards traffic for the user to his real point of presence in the network, the so-called foreign agent. The foreign agent establishes a tunnel to the home agent, using its own IP address as a care-of address for the roaming user. All traffic for the client IP address in the home network can now be forwarded in a tunnel to the foreign agents IP address. The client itself is bound to an address of the subnet at his point of presence.

Advantages

- No dynamical NAT configuration necessary
- L2 connectivity to the home network (e.g. printers)

Disadvantages

- Home- and foreign agents for every network necessary (costly)
- traffic is encapsulated (overhead)

Aruba IP Mobility The wireless network equipment vendor Aruba offers a Mobile IP implementation for his wireless equipment [?, p. 659]. The wireless controllers are used as agents. However, relying only on this commercial solution leads to several different disadvantages:

- Roaming is only available in buildings with a WLAN controller
- Each residential home needs a WLAN controller (costly)
- It is unknown if the solution is technical applicable to our environment
- The whole setup is locked down to Aruba

The technical issues will be investigated in the future. The reason it may not work in our environment is due to the determination of the home agent in the network. According to the documentation the calculation is based on the user VLAN IP subnet [?, p. 661]. To work properly in our network, VLAN ids may not be unique in the whole network.

2.2 Netfilter Hooks

The netfilter hooks[?] allow to tap into the network traffic on a linux machine in specific places in the network stack. The different hooks allow to access the network traffic in specific stages, making filtering, altering

and analysing pretty flexible. Figure 2.2 shows an overview of the different hooks in the network stack. The important hooks for NAT are PREROUTING and POSTROUTING. For source NAT the destination IP address has to be correct in order to change the source IP address correctly. For destination NAT it is the other way round. It must be performed before the routing, otherwise its not clear what the actual destination IP address is and how the packet has to be routed.

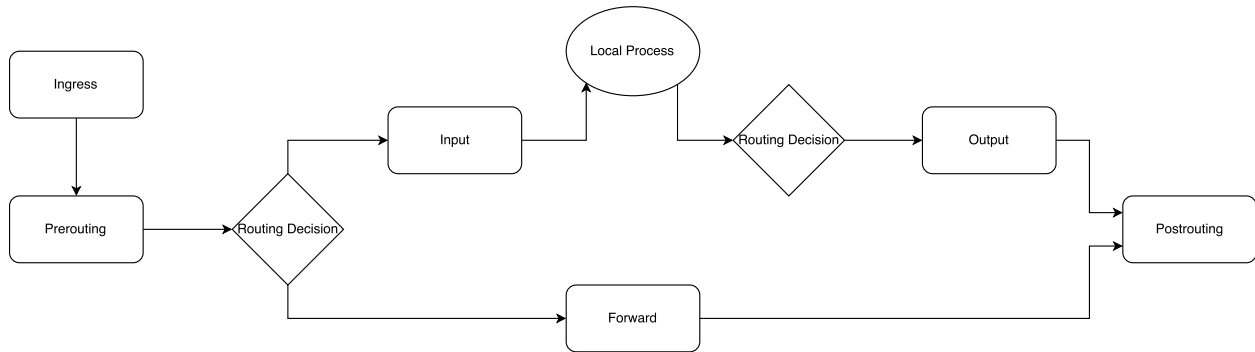


Figure 2.1: netfilter hooks

Chapter 3

Management Component

3.1 Structure

The management component is a python3 celery application. Celery is a framework designed for asynchronous task queue execution[?]. It enables management components in the network to trigger configuration changes at the gateway. Updates are distributed by a rabbitmq message broker[?]. An example configuration for the management component can be found in the appendix at 7.5.

3.2 Data Model

3.2.1 Translations

Translations are used to configure NAT mappings. They contain the public IP address and the private subnet. Private subnets do not have a fixed size. This design allows single forwarding entries for roaming users represented by private subnets as small as a single host (/32). Home networks with an expected size of 254 hosts (/24) can be represented in these objects, too.

3.2.2 Forwardings

These objects are used for DNAT to allow incoming connections to users. They contain the following data: public IP address, private IP address, protocol, source port and destination port.

3.2.3 Throttles

Throttle objects are used to describe limited connections. They contain the private network or host, the public IP address and the limiting speed in kbytes/second. Public and private addresses are both required, because inbound and outbound connections have to be limited. The speed limit is shared for inbound and outbound connections as the same queue is used.

3.3 Program State

The application supports two ways to update its configuration.

First, a task named initialization is possible. It is called after the start of the application. During the initialization, the first configured database is queried to generate a complete ruleset. The conntrack state of the NAT gateway for the configured private network range is dropped. Nftables configuration on all relevant tables is dropped and the new state is applied. The initialization process can later be utilized to reset the gateway.

Second, incremental state updates are possible. They are triggered by the message broker with parameters to exactly identify the required change. During an update, the relevant database is queried for the data. The content of the change is not communicated via the message broker but the database, so the broker only signals the application that an update is required.

3.4 State Updates

As mentioned above, the application queries the database to determinate which action to perform during a signalized state update. Two scenarios are possible: the data may or may not be present in the database. If there is no matched data in the database, the application assumes that the represented data (translation, forwarding or throttle) has to be dropped from nftables configuration. Then it performs the necessary tasks to delete the entries. This may include dropping conntrack state information. If matched data is inside the Tabaksdose, the application rewrites all affected nftables rules. This may result in a newly created rule or in an updated rule, depending on the object in question and the current nftables state. For example, translations allow direct rule replacements. An already present rule may be updated with the “nft replace rule” command.

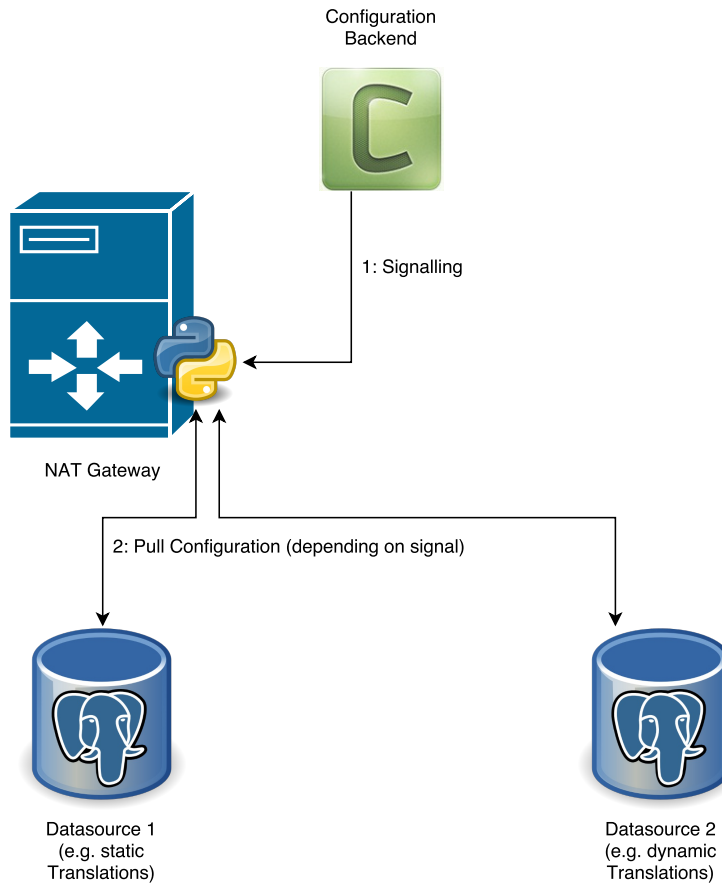


Figure 3.1: State updates

Chapter 4

Implementation

4.1 packet filter

Packet filter (pf)[?][?] is a firewall for the BSD-Family and was original written for OpenBSD. Besides filtering pf is moreover able to perform NAT and bandwidth control. Due to time issues we did not implement a solution with pf, but evaluated some basic constrains. Instead of using OpenBSD we would recommend using FreeBSD, because OpenBSDs PF is not able to utilitise multi core processors [?]. This may causes performance issues, so this feature may be important. Also the Solarflare network cards of our gateways are only supported on FreeBSD. Like the other solutions pf keeps track of connections and only traverses the NAT rule set if the packet does not belong to an existing connection. For redundancy pf uses pfsync[?], which allows the synchronization of state table entries between multiple pf instances.

4.2 pfSense

An related project is pfSense[?], a firewall distribution based on OpenBSD and PF. The software includes an webinterface, is actively maintained and has an online community. Sadly, it's not possible for us to use this existing project, since it does not provide an API to do remote configuration. This feature is planned for the 3rd release of pfSense, but the release date is unknown until further notice.

4.3 Iptables

iptables[?] is the de facto standard for network packet manipulation on Linux. It is a user-space application that allows configuration of the packet filter ruleset in the Linux kernel. Besides filtering, it can be used to perform NAT. There are different implementations of iptables for different protocols like IPv6 and ARP. Iptables has 5 standard tables: filter, nat, mangle, raw and security. Filter is the default table and is used for filtering unwanted packets, it has the INPUT, FORWARD and OUTPUT hooks registered. The nat table is for performing NAT, packets pass this table only when a new connection is created. Its built-in hooks are PREROUTING, OUTPUT and POSTROUTING. To alter packets, the mangle table is used. Raw is for configuring exemptions to the connection tracking and security is used for Mandatory Access Control. Besides the standard tables, new tables can be added to perform specific tasks. Tables contain a set of chains containing a set of rules. Rules match traffic and then perform a specific task, like NAT or jumping to another chain.

4.4 Nftables

nftables[?][?][?] is a new packet classification framework for Linux that aims to replace iptables. It was first introduced on the netfilter workshop 2008 and got its first official release in 2009. Besides iptables, the tools

ip6tables, ebtables and arptables are going to be united in nftables.

The motivation behind this new tool is the age of iptables and many problems like the representation of the ruleset as huge blob, which makes replacement of single rules impossible. The ruleset is not very memory efficient and can not be translated back into rules. Another mayor drawback is that iptables, ip6tables, ebtables and arptables share a large codebase but are not the same tool, which makes the code management inefficient. Nftables also comes with an easier syntax and allows multiple targets per rule, making rulesets less redundant. It is also possible to exchange single rules in a ruleset without reloading the entire table.

nftables consists of three main components: the kernel modules, nft (the nftables user-space application) and libl the netfilter netlink library for communication with the kernel.

The basic container in nftables is a table, it contains rules and chains of the same protocol family. Chains are the containers for rules and can be used as jump targets. A chain with a netfilter hook is called a base chain and serves as entry point into a table. Rules are container of expressions that define the runtime behavior and are the smallest unit that can be replaced. Besides these basic building blocks, nftables allows the usage of sets and maps. They make it easy to specify rules without much redundancy working for the entire used data structure. Especially maps are a useful addition for NAT, because they allow to directly specify NAT mappings. There are two different implementations for sets, rbtree or hashmap. Nftables internally decides which one its going to use.

4.5 Conntrack / Conntrackd

Conntrack[?] is a user-space application that provides access to the connection tracking module[?] of the Linux kernel. The sole use of the connection tracking module is to store information about connections that pass the machine. Each connection is identified by the source and destination IP addresses, the port numbers, and protocol type. The connection tracking module stores information about the state of the connection and is vital for NAT, because it allows to match incoming traffic from the internet to internal hosts.

After the initial NAT through nftables or iptables, a conntrack entry gets created in the connection table, containing the new source IP address. All following packets of the same connection are directly translated using the conntrack entry and do not pass through the nat table.

Example conntrack:

```
icmp 1 4 src=100.64.0.10 dst=192.168.30.0 type=8 code=0 id=7224
src=192.168.30.0 dst=192.168.0.19 type=0 code=0 id=7224 mark=0 use=1
```

The first src, dst matches the original traffic from the internal host, the second pair matches the response from the host on the other side of the gateway. The conntrack contains all the information needed to perform the NAT.

Conntracks can have different states:

- NEW: The connection is starting, to reach this the state the packets only have to be valid, no reply packets have been received yet.
- ESTABLISHED: The connection is established, this means the gateway has seen replies to the connection.
- RELATED: An expected connection, that is used for some special protocols like FTP. They are created by so called helpers that can extract information out of a connection that lets you expect another connection in the future. For example, the control flow for FTP is done over port 20, but the data connection is done with a different port that is sent to the client by the server. The ftp helper extracts that port number and sets up an expected connection for that port, allowing this connection to pass the NAT gateway.
- INVALID: The connection does not follow valid network communication.

The Kernel module uses a hash table for efficient lookups, that holds 2 hash values for every connection: one for the original direction and another for the reply direction. The hash contains all necessary information to match the connection like the source and destination IP address. The two hash values point to the conntrack that holds the state information of the connection. Current connections are not altered during NAT rule updates. Only new connections are immediately affected by the rule changes. To change all connections to a new IP address the old conntrack state needs to be dropped.

Besides the kernel module for conntrack and the user-space tool, there is the daemon conntrackd. The daemon allows state table synchronization between different NAT gateways. Utilising the state synchronization, it is possible to setup fault tolerant NAT gateways that do not lose connections in the case of an failure. In our setup an extra link between the gateways is used to replicate state information with a low-latency. This allows us to use asynchronous routing with our gateways, meaning that a connection can be translated by one gateway in one direction and translated back by the other gateway in the opposite direction. We performed tests that suggest that this setup works, but we have to wait for live tests to be sure. If it does not work, we have to make sure that a connection passes the same gateway in both directions (for example by splitting up the IP addresses for the gateways).

The maximum size of the connection tracking table can be adjusted with the variable *net.ipv4.netfilter.ip_conntrack_max* and should be increased from the default value. Conntrackd also has an option specifying the maximum number of conntracks.

4.6 Keepalived

Keepalived[?] is a Virtual Router Redundancy Protocol (VRRP)[?] daemon for Linux. Its main goal is providing simple loadbalancing and high-availability for routers and servers running Linux. VRRP allows to configure one virtual router with an IP address using a set of hardware routers. The virtual IP address gets assigned to a specific hardware router by VRRP, making this router the VRRP master. The other routers function as backup routers for the master. If the master fails, the VRRP assigns the IP address to a new hardware router. This functionality allows easy fail over for routers and servers.

In our setup both gateways are active at the same time, making it necessary to configure two virtual gateways, each gateway providing backup functionality for the other. We added the configuration file we used for our evaluation scenario in the appendix at 7.4.3.

Chapter 5

Setup

5.1 Live-Setup

The central router in our network is called Ianus and connects to our uplink. Ianus consists of two hardware routers located at different buildings. To ensure availability, one NAT gateway will be located at each of those locations. Each gateway will be connected to both hardware routers of Ianus. This ensures availability in case that one gateway or one hardware router fails.

All outgoing user traffic will be sent to the gateways and from the gateways back to Ianus. Through this loop setup it is very easy to not route traffic over the gateways that should not be translated (e.g. connections to and from our servers).

5.2 Gateways

The general hardware/software specifications for our two NAT gateways (Dell PowerEdge R630 Server):

CPU: 2x Intel Xeon E5-2690 v4 (2,6GHz, 14 Cores/28 Threads)

RAM: 2x 16GB RDIMM 2.400MT/s

first NIC: Intel Ethernet i350 4 Ports 1GbE

second NIC: 2x Solarflare Flareon Ultra SFN8522 10 Gigabit SFP+

harddrive: 200GB SSD

OS: Debian stretch, Kernel ≥ 4.8

nft: > 0.7 (build from the git repo)




Figure 5.1: The live-setup with the connections between the gateways and Ianus




Figure 5.2: This is the normal traffic flow for an internal host that communicates with the internet

Chapter 6

Evaluation

6.1 Introduction

To check if the NAT gateway is capable of serving all our users, we relied on performance/stress testing to see what the gateway is capable of. Generating adequate amounts of network traffic for testing was challenging. Traffic can vary on different characteristics: used protocols, bandwidth, packet amount, packet length and connections.

A simple test with iPerf[?] revealed that pure bandwidth had little to no effect on the NAT gateway. Another test with a few million packets per second showed the same. Only when we started using more connections (different source ip address destination ip address pairs), the cpu usage on the NAT gateway went up. This is because only the first packet of a network connection traverses the NAT rules in nftables or iptables. All following packets are directly processed by the connection tracking system of the kernel. Based on this conclusion, our stress testing relies on creating as many connections as possible.

The problem with connections is that they can not be generated as easy as lots of packets or bandwidth. Mostly because there needs to be an response to a request that traverses the NAT gateway in the opposite direction. This makes techniques like IP-spoofing impractical. All IP addresses in use need to be configured on the corresponding machines (now is a good time to increase the size of your ARP cache).

6.2 Setup

In Figure 6.2 you can see our setup. We used one virtual machine (VM) to simulate hosts in “The Internet” (192.168.0.0/16) and one for our home networks in the Carrier-grade NAT (CGN)(100.64.0.0/12). For “The Internet”, the NAT gateways have all public IP addresses configured that are used for NAT (of course not both gateways at the same time). In the Carrier-grade NAT each gateway only needs one IP. Between the two gateways there is a link for state synchronization of the conntrack tables.

An internal host, simulated by the “Home Networks” VM, sends a packet over one of the gateways to an internet host. After the first packet has traversed the NAT rules, a conntrack is created and synchronized with the second gateway. The internet host receives the packet and sends a response via the gateway that has the corresponding IP address configured. The gateway then forwards the message to the internal host.

6.3 Generating Traffic

After some initial testing with packet generation tools and python, we had problems generating network traffic that puts much load on the machines. The key to generating good workload is to generate lots of different connections, meaning lots of different source IP and destination IP pairs. To generate them easily we found the unix tool nping[?] very helpful. Nping is part of nmap[?] and supports multiple target IPs (up to 8900) and gives you a lot of control about how the packets are generated. It supports TCP, UDP and

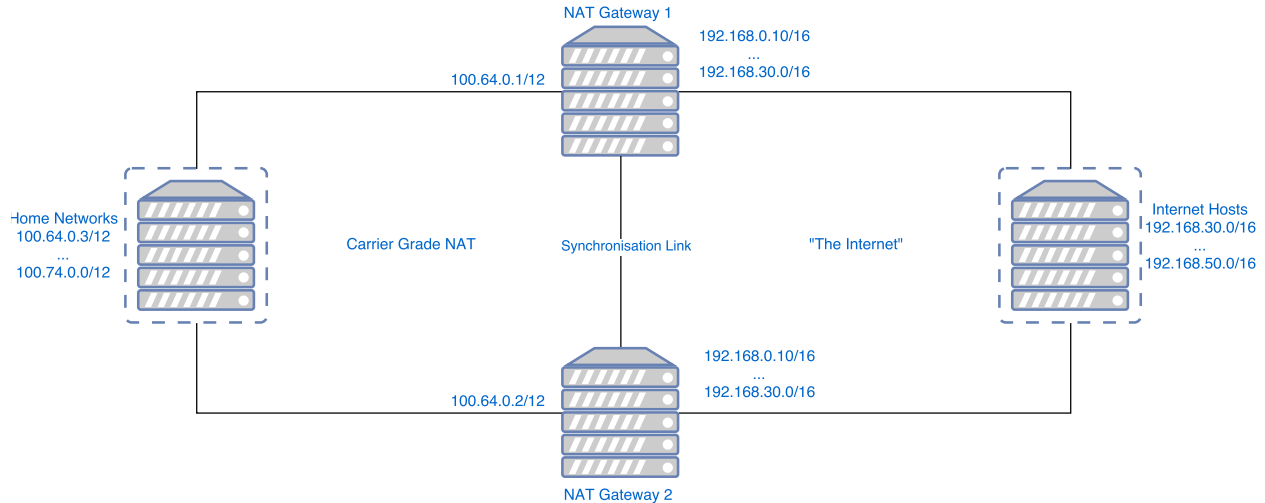


Figure 6.1: Evaluation setup

ICMP packet generation and can use a specified source IP address, port and interface. The payload of the nping packets can be specified as well. This feature set makes nping an excellent tool for our use case.

To generate more connections, we started multiple instances of nping with different source IP addresses. We will include the used python script in the appendix.

6.4 Metrics

The metrics we are interested in are latency, resource consumption on the gateway and the synchronization of conntracks. All of them together should help us to find bottle necks and identify problems in the gateways.

The latency is interesting because it should be as low as possible to not harm the hosts that are behind the NAT. We measured the latency by using ping[?] from the internal hosts VM to the Internet Hosts VM. To measure the resource consumption we used collectl[?], a unix tool that can measure various resources on a pc. We used it to measure the cpu utilization, network traffic and interrupts on the gateways. The polling interval was one second. For the synchronization of the state tables we used “conntrackd -s” and polled this statistic every second with a script. Another metric we looked at was the memory consumption, but we could not find anything interesting here, so we excluded it.

6.5 Problems

While working on the evaluation and the traffic generation we encountered two mayor bugs: When the conntrack table is full conntracks get dropped, this works fine for small tables but results in stuck CPUs with bigger ones. We do not know the exact size for the conntrack table to trigger this bug, but it already happend with the default size of 260 000. This bug is severe because it allows to perform very easy denial of service attacks against the NAT gateway. The needed bandwidth and packets per second (pps) to perform this attack are below 20 Mbit/s and 40k pps. We will submit a bug report to the netfilter project, but so far we do not know how to fix this (besides using a small table, which sucks).

The second problem is the “disappearance” of the Solarflare NIC interface that receives/sends the NAT traffic. The NIC interface gets unresponsive and can only be brought up again by a reboot. NICs do so called flow steering which matches connections to the correct core of the CPU to handle the traffic. It could be that this feature plus the high number of connections are the cause of this, but this is only speculation.

6.6 Results

Coming soon to an email server of your choice. . .

Chapter 7

Appendix

7.1 keepalived.conf

```
global_defs {
    # Name of the local router
    router_id      chomsky

    # The email server
    smtp_server    192.168.0.30

    # Where to send notifications if a machine is down
    notification_email {
        test@test.de
    }
}

# Groups of vrrp_instances that failover together, allows to call
# scripts and send notifications
vrrp_sync_group agdsn_nat_1 {
    group {
        cgn_1
        ext_1
    }

    smtp_alert
    global_tracking
}

# Describes moveable IPs, sets the priority for the machine
vrrp_instance cgn_1 {
    state          BACKUP

    interface      eth7

    virtual_router_id 10

    priority        150

    virtual_ipaddress {
        100.64.0.1/12
    }
}

vrrp_instance cgn_2 {
    state          BACKUP

    interface      eth7

    virtual_router_id 11
```

```

        priority      100
        virtual_ipaddress {
            100.64.0.2/12
        }
    }

    vrrp_instance ext_1 {
        state          BACKUP

        interface      eth6

        virtual_router_id 12

        priority        150

        virtual_ipaddress {
            192.168.0.1/16
        }
    }

    vrrp_instance ext_2 {
        state          BACKUP

        interface      eth6

        virtual_router_id 13

        priority        100

        virtual_ipaddress {
            192.168.0.2/16
        }
    }
}

```

7.2 nping

7.2.1 Example usage

Example usage:

```
nping -c 0 --tcp -p 80-1080 --interface eth1 -S 100.64.0.10 -g 5000 --rate 6000 <list of dest ips>
```

```

-c                -> Number of repeats, 0 is for infinity
--tcp            -> Probe mode
-p 80-1080        -> destination port, or port range
--interface eth1 -> the interface to use
-src             -> the source ip address of the interface to use (if there are multiple)
-g              -> the src port to use
--rate 6000      -> the packets per minute

```

7.2.2 Python script

```
import ipaddress
from subprocess import call
```

```
IP = ipaddress.ip_address('192.168.30.0')
```

```

dest_ips = ""
for i in range(8900):
    dest_ips += str(IP) + " "
    IP += 1

```

```
SRC_IP = ipaddress.ip_address("100.64.0.10")
```

```
call("date", shell=True)
```

```

for i in range(30):
    call("nping -c 0 --tcp --interface eth1 -S " + str(SRC_IP) + " -g "
        + str(5000 + i) + " --rate 6000 " + dest_ips + ">> /dev/null &", shell=True)
    SRC_IP += 256

```

7.3 nftables

7.3.1 map

This is a simple Map config for nftables. We wanted to use Maps for our solution but found that they are still too buggy to be used in a live setup. We hope to utilize them in the future, since they are simpler than building a tree. Maps use a RBTree internally so the complexity is still $\log(n)$ like in our other solution. The issues we had so far with maps:

- Adding new elements to a map had a memory leak (this was fixed by the maintainer)
- Very high memory consumption when updating a map entry
- Only the first two rules in a map worked (the maintainer sent us a patch for the kernel module but we still have to try it)

```

#!/usr/sbin/nft
add chain nat postrouting { type nat hook postrouting priority 100 ;}
add chain nat prerouting { type nat hook prerouting priority 0 ;}
add map nat subnettoip { type ipv4_addr: ipv4_addr ; flags interval ; }
add rule ip nat postrouting snat ip saddr map @subnettoip;
add element nat subnettoip { 100.64.0.0/24 : 192.168.0.19 }
add element nat subnettoip { 100.64.1.0/24 : 192.168.0.20 }
add element nat subnettoip { 100.64.2.0/24 : 192.168.0.21 }
add element nat subnettoip { 100.64.3.0/24 : 192.168.0.22 }

```

7.3.2 tree

We have up to 6000 /24 Networks that we want to NAT to specific IP addresses. Simply adding them all in one chain would be too slow, that is why we created a tree of chains. Starting with 2 networks at the top and then 8 networks per level. In the leaf chain the actual NAT rule is defined. Below is the start of our tree config for nftables down to the first 2 blocks of NAT rules. The rest of the tree is structured the same.

```

#!/usr/sbin/nft
add table nat
add chain nat prerouting { type nat hook prerouting priority 0 ;}
add chain nat postrouting { type nat hook postrouting priority 0 ;}
add chain nat postrouting-level-0
add rule nat postrouting ip saddr 100.64.0.0/12 goto postrouting-level-0
add chain nat postrouting-level-0-0
add rule nat postrouting-level-0 ip saddr 100.64.0.0/15 goto postrouting-level-0-0
add chain nat postrouting-level-0-0-0
add rule nat postrouting-level-0-0 ip saddr 100.64.0.0/18 goto postrouting-level-0-0-0
add chain nat postrouting-level-0-0-0-0
add rule nat postrouting-level-0-0-0 ip saddr 100.64.0.0/21 goto postrouting-level-0-0-0-0
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.0.0/24 snat 192.168.0.19
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.1.0/24 snat 192.168.0.20
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.2.0/24 snat 192.168.0.21
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.3.0/24 snat 192.168.0.22
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.4.0/24 snat 192.168.0.23
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.5.0/24 snat 192.168.0.24
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.6.0/24 snat 192.168.0.25
add rule nat postrouting-level-0-0-0-0 ip saddr 100.64.7.0/24 snat 192.168.0.26
add chain nat postrouting-level-0-0-0-1
add rule nat postrouting-level-0-0-0 ip saddr 100.64.8.0/21 goto postrouting-level-0-0-0-1
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.8.0/24 snat 192.168.0.27
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.9.0/24 snat 192.168.0.28
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.10.0/24 snat 192.168.0.29
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.11.0/24 snat 192.168.0.30

```

```
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.12.0/24 snat 192.168.0.31
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.13.0/24 snat 192.168.0.32
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.14.0/24 snat 192.168.0.33
add rule nat postrouting-level-0-0-0-1 ip saddr 100.64.15.0/24 snat 192.168.0.34
```

7.3.3 HowTo

This is a short collection of useful nftables commands we used during our work.

Useful to know:

- If you use the type “nat” for a hook than only the first packet of each connection is processed by the chain
- For NAT there always have to be chains with the prerouting and postrouting hook. For source NAT the prerouting chain is empty but still needed!
- Jump returns to the calling chain after completion, goto never returns.
- If a rate limit rule gets exhausted, the next rule in the chain is called, or the default policy applies. If a host should be throttled, all traffic that exhausts the rate limit needs to be dropped.

```
# Adds the table nat to nftables
nft add table nat

# List the content of the table nat
nft list table nat

# List all tables in nftables
nft list tables

# Can also be used to make a backup of the ruleset
nft list table nat > backup.ruleset

# Load the ruleset backup.ruleset into nftables, see the tree config for an example config
nft -f backup.ruleset

# Deletes the table nat
nft delete table nat

# Deletes all rules in the table nat, chains are preserved
nft flush table nat

# Adds the chain prerouting to the table nat. The chain has the type nat
# and uses the prerouting hook with priority 0. Priority 0 means the its first.
nft add chain nat prerouting { type nat hook prerouting priority 0 \; }

# Add the chain postrouting to the table nat
nft add chain nat postrouting { type nat hook postrouting priority 100 \; }

# Deletes the chain postrouting in the table nat
nft delete chain nat postrouting

# Add a rule to the postrouting chain in the table nat. If the source ip address
# is in the network 100.64.0.0/15 goto chain targetChain
add rule nat postrouting ip saddr 100.64.0.0/15 goto targetChain

# Adds a snat rule to the chain postrouting in the table nat. All source ips from
# 100.95.255.0/24 of the interface eth0 get translated to 141.30.233.255
nft add rule nat postrouting ip saddr 100.95.255.0/24 oif eth0 snat 141.30.233.255

# Addes a rule to the input chain in the table filter. Rate limits the traffic
# to 10 mbytes/second and accepts the traffic. When the rate limit is exhausted
# the next rule is applied. If you want to throttle a host then a drop rule is needed.
nft add rule filter input limit rate 10 mbytes/second accept
```

```

# Adds a new named map subnettoip to the table nat. The map projects ipv4 addresses
# to ipv4 addresses. The interval flag allows you to use not just single addresses
# but entire networks: { 100.64.0.0/10 : 192.168.0.1 }
nft add map nat subnettoip { type ipv4_addr; flags interval \; }

# Adds an element to the map subnettoip in the table nat.
nft add element nat subnettoip { 100.64.1.0/24 : 141.30.202.171 }

# Deletes the element 1.2.3.0/24 from the map subnettoip
nft delete element nat subnettoip { 1.2.3.0/24 }

# Lists the rule handles for the table nat
nft list table nat -a

# Deletes the rule with the handle 4
nft delete rule nat prerouting handle 4

# Adds the a snat rule to the postrouting chain in the table nat. The snat is
# performed using the subnettoip map.
nft add rule ip nat postrouting snat ip saddr map @subnettoip;

```

7.4 iptables

7.4.1 tree

This is the equivalent to the nftables tree for iptables.

```

*nat
:PREROUTING ACCEPT
:INPUT ACCEPT
:POSTROUTING ACCEPT
:postrouting-level-0 -
-A POSTROUTING -s 100.64.0.0/12 -g postrouting-level-0
:postrouting-level-0-0 -
-A postrouting-level-0 -s 100.64.0.0/15 -g postrouting-level-0-0
:postrouting-level-0-0-0 -
-A postrouting-level-0-0 -s 100.64.0.0/18 -g postrouting-level-0-0-0
:postrouting-level-0-0-0-0 -
-A postrouting-level-0-0-0 -s 100.64.0.0/21 -g postrouting-level-0-0-0-0
-A postrouting-level-0-0-0-0 -s 100.64.0.0/24 -j SNAT --to 192.168.0.19
-A postrouting-level-0-0-0-0 -s 100.64.1.0/24 -j SNAT --to 192.168.0.20
-A postrouting-level-0-0-0-0 -s 100.64.2.0/24 -j SNAT --to 192.168.0.21
-A postrouting-level-0-0-0-0 -s 100.64.3.0/24 -j SNAT --to 192.168.0.22
-A postrouting-level-0-0-0-0 -s 100.64.4.0/24 -j SNAT --to 192.168.0.23
-A postrouting-level-0-0-0-0 -s 100.64.5.0/24 -j SNAT --to 192.168.0.24
-A postrouting-level-0-0-0-0 -s 100.64.6.0/24 -j SNAT --to 192.168.0.25
-A postrouting-level-0-0-0-0 -s 100.64.7.0/24 -j SNAT --to 192.168.0.26
:postrouting-level-0-0-0-1 -
-A postrouting-level-0-0-0 -s 100.64.8.0/21 -g postrouting-level-0-0-0-1
-A postrouting-level-0-0-0-1 -s 100.64.8.0/24 -j SNAT --to 192.168.0.27
-A postrouting-level-0-0-0-1 -s 100.64.9.0/24 -j SNAT --to 192.168.0.28
-A postrouting-level-0-0-0-1 -s 100.64.10.0/24 -j SNAT --to 192.168.0.29
-A postrouting-level-0-0-0-1 -s 100.64.11.0/24 -j SNAT --to 192.168.0.30
-A postrouting-level-0-0-0-1 -s 100.64.12.0/24 -j SNAT --to 192.168.0.31
-A postrouting-level-0-0-0-1 -s 100.64.13.0/24 -j SNAT --to 192.168.0.32
-A postrouting-level-0-0-0-1 -s 100.64.14.0/24 -j SNAT --to 192.168.0.33
-A postrouting-level-0-0-0-1 -s 100.64.15.0/24 -j SNAT --to 192.168.0.34

```

7.4.2 HowTo

```

# Show all chains of the table nat
iptables -t nat -L

# Add the chain custom-chain to the table nat
iptables -t nat -N custom-chain

```



```

# Removes the chain custom-chain
iptables -t nat -X custom-chain

# Add the rule to the chain POSTROUTING in the table nat.
# The rule jumps to the chain custom-chain if the source ip address matches.
iptables -t nat -A POSTROUTING -s 100.64.0.0/12 -j custom-chain

# Removes the jump rule from the chain POSTROUTING
iptables -t nat -D POSTROUTING -j custom-cahin

# Change the source ip to 123.123.123.123
iptables [...] -j SNAT --to-source 123.123.123.123

# Uses masquerade on the source ip address -> changes it to the ip address of the outgoing interface
iptables [...] -j MASQUERADE

# Changes the destination ip address and port to 123.123.123.123:22
iptables [...] -j DNAT --to-destination 123.123.123.123:22

# Saves the ruleset
iptables-save > dump.ruleset

# Loads the ruleset
iptables-restore < dump.ruleset

```

7.4.3 conntrackd.conf

```

# Sync section to specifies how the synchronisation is done,
# which protocol is used, etc.
Sync {
    # The synchronisation mode and its options
    Mode FTFW {
        ResendQueueSize 131072
        PurgeTimeout 60
        ACKWindowSize 300
        # Disables the external cache, the state entries
        # are directly injected in the state table
        DisableExternalCache On
    }
    # The protocol that is used for synchronisation and
    # the used ip addresses, buffer size etc.
    UDP {
        IPv4_address 10.77.0.1
        IPv4_Destination_Address 10.77.0.2
        Port 3781
        Interface net1
        SndSocketBuffer 1249280
        RcvSocketBuffer 1249280
        Checksum on
    }
    # General options
    Options {
        TCPWindowTracking On
        # Also sync the expectation table
        ExpectationSync On
    }
}

# General options
General {
    #Systemd support
    Systemd on
    Nice -20
    Scheduler {
        Type FIFO
        Priority 99
    }
    # Number of buckets in the cache hashtable
    HashSize 32768
    # Number of conntrack entries, should be twice the size

```

```

    # of /proc/sys/net/netfilter/nf_conntrack_max
    HashLimit 131072
    LogFile on
    Syslog off
    LockFile /var/lock/conntrack.lock
    UNIX {
        Path /var/run/conntrackd.ctl
        Backlog 20
    }
    NetlinkBufferSize 2097152
    NetlinkBufferSizeMaxGrowth 8388608
    NetlinkOverrunResync On
    NetlinkEventsReliable Off
    EventIterationLimit 100
}

```

7.5 Application config

To launch the application, checkout the code and start an celery worker at the project root:

```
celery -A tasks worker --loglevel=info
```

The configuration file dsnat.cfg:

```

netfilter:
  tree:
    # The cidr size of the last tree element
    lowlevel: 21
    # Maximum rules on a tree level
    jumpcount: 8
  nft:
    call: '/usr/local/sbin/nft'
    tmpfile: '/tmp/rules.nft'
  forwarding:
    table: nat
  translation:
    table: nat
  throttle:
    table: filter
    map: throttle_map

conntrack:
  call: '/usr/sbin/conntrack'

# Multiple databases can be configured
databases:
  - name: static
    host: 10.10.233.1
    user: nat
    password: test123
    db: nat-static
  - name: dynamic
    host: 10.10.233.1
    user: nat
    password: test123
    db: nat-dynamic

broker:
  host: 10.10.233.1
  user: natgateway
  password: test123
  queue: nat

cgn:
  ip: 100.64.0.1
  net: 100.64.0.0/11
  interface: eth4

```

```
inet:
  interface: eth7

# connections from/to these networks are not throttled
whitelist:
  - 141.30.0.0/16
  - 141.76.0.0/16

# these are hosts in the whitelist networks, that should not ne spared
blacklist:
  - 141.30.3.50
  - 141.30.4.125
  - 141.30.61.140
  - 141.30.61.141
  - 141.30.117.36
```